

**Task 1a. Look at the ls code in user/ls.c and explain how the code works in terms of the library and system calls it makes. Show the output from the ls command and interpret the different columns.**

The ls code in user/ls.c implements the command utility ls that is used to list the contents of the directory. The code includes header files like types.h, stat.h and fs.h from kernel and user.h from user. These files provide definitions and declarations for types, file systems and user-level functions. The first function 'fmtname' formats the name of a file and pads the name with spaces to ensure that it has more than the fixed length 'DIRSIZ+1' and returns the newly formatted name. The second function 'ls' takes the path as an argument and list the contents of the directories of that specific path. It opens the directory using 'open' and retrieves the information using 'fstat'. The third 'main' function is where it calls the 'ls' function for the current dir if arg<2 or else it calls ls for the arg provided.

```
sstshering@DESKTOP-20BNCM1:~/OS-riscv-xv6/user$ ls
_cat      cat.sym      initcode.asm mkdir.o      sh.d         uptime.c
_echo     echo.asm     initcode.d   mkdir.sym    sh.o         uptime.d
_forktest echo.c        initcode.o   printf.c     sh.sym       uptime.o
_grep     echo.d       initcode.out printf.d      sleep.asm    uptime.sym
_grind    echo.o       kill.asm     printf.o     sleep.c      user.h
_init     echo.sym     kill.c       ps.asm       sleep.d      usertests.asm
_kill     forktest.asm kill.d        ps.c         sleep.o      usertests.c
_ln       forktest.c   kill.o       ps.d         sleep.sym    usertests.d
_ls       forktest.d   kill.sym     ps.o         stressfs.asm usertests.o
_matmul   forktest.o   ln.asm       ps.sym       stressfs.c   usertests.sym
_mkdir    grep.asm     ln.c         ptest.asm    stressfs.d   usys.S
_ps       grep.c       ln.d         ptest.c      stressfs.o   usys.d
_ptest    grep.d       ln.o         ptest.d      stressfs.sym usys.o
_pstree   grep.o       ln.sym       ptest.o      time.d       usys.pl
_rm       grep.sym     ls.asm       ptest.sym    time1.asm    wc.asm
_sh       grind.asm    ls.c         pstree.asm   time1.c      wc.c
_sleep    grind.c      ls.d         pstree.c     time1.c.save wc.d
_stressfs grind.d       ls.o         pstree.d     time1.d      wc.o
_time1    grind.o      ls.sym       pstree.o     time1.o      wc.sym
_uptime   grind.sym    matmul.asm   pstree.sym   time1.sym    zombie.asm
_usertests init.asm     matmul.c     rm.asm       ulib.c       zombie.c
_wc       init.c       matmul.d     rm.c         ulib.d       zombie.d
_zombie   init.d       matmul.o     rm.d         ulib.o       zombie.o
cat.asm   init.o       matmul.sym   rm.o         umalloc.c    zombie.sym
cat.c     init.sym     mkdir.asm    rm.sym       umalloc.d
cat.d     initcode     mkdir.c      sh.asm       umalloc.o
cat.o     initcode.S   mkdir.d      sh.c         uptime.asm
```

**Task 1b. Look at the mkdir code in user/mkdir.c and explain how the code works in terms of the system call it makes and the OS helper functions used by that system call. Use mkdir and some other commands (e.g., echo) to create a new directory and a couple of files in it. Run ls again and explain the output.**

Mkdir.c also has header files to provide definitions and declarations for types, file systems and user-level functions. The main function being the entry point of the code, it checks if the number of command line arguments, or argc < 2 which means no directory names were given. In that case, it prints out a message and exits the code with 0. Then it iterates through the provided

directory names, `argv[i]` and tries to create each directory using **mkdir** (sys call used to create new directory) system call. After creating directories successfully, it exits with 1.

```
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6$ mkdir my-dir
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6$ ls
LICENSE  README  kernel  my-dir
Makefile fs.img  mkfs    user
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6$ ls -l
total 1032
-rw-r--r-- 1 sstshering sstshering  1174 Sep 19 13:04 LICENSE
-rw-r--r-- 1 sstshering sstshering  4793 Dec 15 11:25 Makefile

-rw-r--r-- 1 sstshering sstshering  2226 Sep 19 13:04 README
-rw-r--r-- 1 sstshering sstshering 1024000 Dec 13 02:28 fs.img
drwxr-xr-x 2 sstshering sstshering  4096 Dec 15 11:25 kernel
drwxr-xr-x 2 sstshering sstshering  4096 Dec 13 03:33 mkfs
drwxr-xr-x 2 sstshering sstshering  4096 Dec 15 12:00 my-dir
drwxr-xr-x 2 sstshering sstshering  4096 Dec 15 11:25 user
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6$ cd my-dir
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6/my-dir$ ls
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6/my-dir$ echo "This is
a new file" > file.txt
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6/my-dir$ ls
file.txt
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6/my-dir$
```

**Task 1c.** Look at the `ln` code in `user/ln.c` and explain how the code works in terms of the library and system calls it makes. Use `ln` to create one or more links to a file you created in Task 1b. Run `ls` again and explain the output.

The `ln.c` also has header files to provide definitions and declarations for types, file systems and user-level functions. The main function checks if `argc!=3`, which means incorrect usage, and prints usage message and exits with error code. Then the **link** (used to create a new hard link to an existing file) system call is used to create a hard link between the old file and the new link, which is depicted by `argv[1]` and `argv[2]`. After the link has been created, the code exits with 0 = success, and if <0, error.

```
sstshering@DESKTOP-20BNM1:~$ ln file.txt linkfile.txt
sstshering@DESKTOP-20BNM1:~$ ls
OS-riscv-xv6  file.txt  linkfile.txt  mydir
sstshering@DESKTOP-20BNM1:~$
```

**Task 2a.** Look at the code in `kernel/sysfile.c` for the `fstat()` system call and in `kernel/file.c` for the `filestat()` helper function and explain how the codes work.

The **fstat()** system call in `sysfile.c` is responsible for retrieving file status information. It fetches the **argfd** (file descriptor) as an argument and checks if its valid and then

fetches **argaddr** (user space addr) for the **struct stat**. It then calls **the filestat()** to retrieve the file info and returns the result of this call, otherwise -1 on error/failure.

The **filestat()** sys call on **file.c** checks if the file type is supported, if so it locks the related inode and retrieves the file info using the **stati** (copies relevant info from inode to the struct stat) function and then unlocks the inode. It then copies the file info to the user space.

**Task 2b. On a Linux system, look at the stat man page. Run the stat command on a regular file and on a directory and explain the output.**

The stat command on file and directory gives output with the file name, size, blocks and all other related information regarding the file/ directory.

```
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6/kernel$ stat bio.c
  File: bio.c
  Size: 3325          Blocks: 8          IO Block: 4096   regu
lar file
Device: 820h/2080d   Inode: 26966          Links: 1
Access: (0644/-rw-r--r--)  Uid: ( 1000/sstshering)   Gid: ( 1000
/sstshering)
Access: 2023-12-13 01:44:11.084203230 -0700
Modify: 2023-09-19 13:04:33.213056861 -0600
Change: 2023-09-19 13:04:33.213056861 -0600
 Birth: 2023-09-19 13:04:33.213056861 -0600
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6/kernel$ cd..
cd..: command not found
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6/kernel$ cd ..
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6$ ls
'CS4375 Fall2023 HW1 Report.pdf'  Makefile  fs.img  mkfs
LICENSE                           README    kernel  user
sstshering@DESKTOP-20BNM1:~/OS-riscv-xv6$ stat user
  File: user
  Size: 4096          Blocks: 8          IO Block: 4096   dire
ctory
Device: 820h/2080d   Inode: 29418          Links: 3
Access: (0755/drwxr-xr-x)  Uid: ( 1000/sstshering)   Gid: ( 1000
/sstshering)
Access: 2023-12-15 12:19:22.865094567 -0700
Modify: 2023-12-15 12:19:21.261761241 -0700
Change: 2023-12-15 12:19:21.261761241 -0700
```

**Task 2c. In your hw7 branch, implement an fstat command for xv6 that takes the name of a regular file or a directory as an argument and outputs the information for that file. The information output should include the filename, file type, file size, inode number, and number of links in an easily understandable format.**

```
uint64
sys_fstat(void)
{
    char path[MAXPATH];
    struct file *f;
    struct stat st;

    if (argstr(0, path, MAXPATH) < 0)
        return -1;

    if (filestat(f, (uint64)&st) < 0)
        return -1;

    printf("File: %s\n", f);
    printf("Type: %s\n", (st.type == T_FILE) ? "Regular File" : "Directory");
    printf("Size: %d bytes\n", st.size);
    printf("Inode: %d\n", st.ino);
    printf("Links: %d\n", st.nlink);

    return 0;
}
```